



# CorpusMind Lens User Guide

A two-page companion for researchers: install the app, build image and social media corpora, run the classic text toolset and the visual battery, and publish reproducible results, entirely on your own machine.

VERSION 0.3.2

SEPTEMBER 2026

ENGLISH + ARABIC INTERFACE

OFFLINE FIRST

**CorpusMind Lens** is a desktop laboratory for multimodal corpus research. It analyses images, posters and social media text in English and Arabic, keeps every file and model on your machine, and never sends your corpus anywhere. Version 0.3 reorganised the interface around the research workflow and added a complete text-analysis layer; version 0.3.2 refines it around the first testing round: honest model pickers, a Save dialog for every export, an explicit local-LLM mode, and research-accurate analysis names. This guide matches version 0.3.2 exactly.

## 1 Install and first run

FIVE MINUTES, NO CONFIGURATION

- **Download one installer** from the release page: setup.exe or msi on Windows, dmg on macOS, deb or ApptImage on Linux. Installers are large because the vision model stack and the reference frequency tables ship inside them.
- **Open Lens and go.** The app starts its own analysis engine automatically; the first launch self-extracts the bundled model stack, so allow it a little extra time once.
- **Take the welcome tour.** A three-page walkthrough explains the workflow on first start; replay it any time from Settings.
- **No port conflicts, ever.** The engine picks the first free port from 8765 upward on its own. If CorpusMind (Text) is already running, Lens detects it and the two can share one conversation through Companion Mode.

**Prefer a browser?** Start the engine, open the web interface, and use the install icon in the address bar to add Lens as a PWA. Language and theme live in the ribbon; Arabic mirrors the whole interface right-to-left.

## 2 Set up local AI models

SETTINGS, THEN AI BACKEND AND MODELS

- **One-click backend.** The card detects Ollama or LM Studio. If neither is present, press Install Ollama: the setup is silent, per-user, and needs no admin password on any platform.
- **Honest fit advice.** Lens reads your RAM, CPU and GPU, then badges every model GPU-fast, CPU OK, Tight or Too big, and recommends sensible defaults for this machine.
- **Pull with live progress.** Filter the catalogue by task (vision, text, embeddings) or search thousands of Hugging Face models; progress streams in real time and models can be deleted just as easily.
- **Choose defaults per feature.** The Default models card lists only models actually downloaded on this machine, merged across Ollama and LM Studio; a default that is missing is flagged and can be reset to the built-in choice in one click.

**No backend installed?** Every deterministic feature still works. A separate Settings card reports each analysis capability honestly, so you always know what is active.

## 3 The menu follows the workflow

FULL FUNCTIONAL MENU, YOUR CHOICE

- **One tool per page, in research order.** The sidebar groups everything the way you work: Start, then 1 Build corpus (Images, Social Media), then 2 Analyse (Text Analysis, Vision Analysis, Assistant), then 3 Share (Export and Methods), then Settings, User Guide and About.
- **The task bar at the bottom always tells the truth:** engine state, how many images are still processing, and a suggested next step for where you are in the workflow.
- **Ctrl or Cmd + K** opens a command palette from anywhere: jump to any page, flip the theme, or switch language; a hint chip in the task bar opens it too.
- **Nothing is hidden.** There are no tabs inside tabs; the set you are working on is selected at the top of every analysis page.

## 4 Create a project and upload images

IMAGES PAGE

- **Projects group your work** (for example, Election posters 2025), created on Home. Inside a project, a New image set is your corpus unit: the batch of images you want to analyse together.
- **Upload on the Images page:** JPG, PNG, TIFF, WebP or BMP. Ingestion is hardened, so a file mislabelled with the wrong extension is rejected with a clear message instead of crashing the batch.
- **Analysis runs in the background,** and the task bar counts every pending image. Thumbnails fill in as each one reaches Ready; open any image to inspect its OCR text, colours, detections and typography.
- **Annotate on the spot.** Tag the five visual dimensions per image, multi-select values, and bulk-tag a whole set as reviewed. Delete, re-run OCR or re-run the full analysis of any image from its always-visible detail header.

## THE INTERFACE AT A GLANCE

### Home

Workflow, projects, engine status.

### Images

Sets, uploads, annotation, gallery.

### Text Analysis

Word list, concordance, keyness.

### Vision Analysis

Measures, grammar, framework lenses.

### Social Media

Archive imports, API fetches, analyses.

### Assistant

Grounded answers, audit log.

### Export & Methods

Tables, corpus files, Methods draft.

### Settings

Models, defaults, ethics, diagnostics.

### User Guide

Ten short chapters, in the app.

### Task bar

Progress, issues, next step, always.

**5 Text Analysis over the OCR corpus**

NEW IN 0.3.0, ANTCONC PARITY

- **Word list** gives frequencies per million with type-token ratio, filtered by the stoplist you choose.
- **Concordance** finds keyword-in-context lines from plain words or regular expressions, with sort positions, per-hit image names, and CSV or XLSX export.
- **Collocations** rank co-occurring words by MI, t-score, log-likelihood, LogDice or delta P, drawn as an interactive network; each word opens a sketch of its visual co-patterns (colours, annotations, objects).
- **N-grams and dispersion** surface repeated bundles and how evenly each word spreads across the set (Juilland's D, Gries' DP).
- **Keyness against bundled references** (BE06 and Leipzig for English, CAMEL and Quranic lists for Arabic) works without a second set; significance and effect size are always reported together.

Your own stoplists apply to every word list, collocation and n-gram run.

**6 Vision Analysis**

MEASURES, GRAMMAR, LENSES

- **The Measures battery** computes frequency, diversity (TTR, Guiraud's R, MATTR, STTR), sequence n-grams and dispersion for any annotated dimension.
- **Visual grammar** reads a single image through Kress and van Leeuwen's three metafunctions with cited evidence and confidence values.
- **Twelve theoretical lenses** (Fairclough, van Dijk, Wodak, Barthes, Peirce and more) return claims phrased as hypotheses, with cited evidence and an honest ungrounded flag; pick Heuristic or Local LLM mode, which runs the chat model you set as default.

**7 The task bar and Smart Troubleshooting**

NEW IN 0.3.0

- **Always visible:** engine state on the left, background processing in the middle, and the suggested next step on the right.
- **When something fails,** a red badge counts the issues: a plain-language explanation, an instant offline fix, and, with a local model installed, a deeper interpretation on your own machine.
- **Nothing goes to the cloud.** Interpretation uses your local Ollama or LM Studio model only; issues can be resolved by hand or auto-resolve when the engine recovers, plus a ready-made report-to-developer email.

**8 Ethics and privacy**

ON BY DEFAULT, IN YOUR FAVOUR

- **Facial and body analysis is opt-in and off by default;** person-descriptive model output is redacted until you enable it.
- **No identity recognition, ever.** Lens never re-identifies real people, and GPS metadata is never extracted under any setting.
- **Cloud AI is off by default** and shows a visible indicator whenever enabled; pseudonymised handles use a salt that stays on your machine.

**9 Build social media corpora**

SOCIAL MEDIA PAGE, TWO WAYS

- **Import, fully offline.** Download your own data from the platform, then drop it in: the X (Twitter) archive zip, Instagram or Facebook Download-Your-Information JSON, TikTok export JSON, or any CSV or JSONL file with a text column; attached photos flow into the vision pipeline automatically.
- **Ethics gate first.** Confirm an attestation, then choose pseudonymised handles and URL or mention redaction; email addresses and phone numbers are always redacted. A provenance record is stored with the corpus.
- **Fetch through official free APIs.** Mastodon needs no key, Reddit uses a free script app, YouTube uses a free API key; credentials are used once and never stored, only public content, within free rate limits.
- **Named for what they are.** Textual Analyses covers word frequency, lexical diversity, n-grams and KWIC; Engagement and Network Analyses covers emoji and hashtag profiles, the hashtag co-occurrence network, engagement statistics, engagement-weighted frequency, the time series and keyness; every run exports, and a note points to CorpusMind for deeper analysis.

**10 The AI Assistant**

GROUNDED IN YOUR DATA

- **Ask in plain language:** analyse this poster with Kress and van Leeuwen, or compare two campaigns' compositional patterns.
- **Answers are grounded.** Every claim comes from a tool call into your actual data, shows the tool that produced it, anything ungrounded is flagged in red, and every tool call is written to a local audit log.

**11 Export and reproducibility**

EXPORT &amp; METHODS PAGE

- **Every result exports** as CSV, TSV, XML, JSON or XLSX from one page: the image-set archive, the battery tables and the social corpus, plus the extracted text as a doc-marked corpus file for AntConc, WordSmith or CorpusMind (Text).
- **Spreadsheets stay safe** against formula injection, so a leading equals sign can never execute in Excel. **Choose where files go:** on the desktop every export opens a Save dialog starting in Downloads, and the toast confirms the exact path; in the browser the toast points at Downloads.
- **The Methods paragraph drafts itself,** naming every model and formula version that actually ran, ready to paste into your manuscript.

**12 Troubleshooting and diagnostics**

MOST ANSWERS LIVE HERE

- **Watch the task bar:** it flags every failed request and shows the engine state; the in-app guide has short chapters on daily tasks.
- **Stuck engine?** Settings has an Engine diagnostics card: restart the engine and read its log tail without leaving the app; logs live in .lens-engine-data under your user profile.
- **AI features unavailable?** Open Settings, AI backend and models, confirm the backend is running and a model is pulled, then pick defaults in the Default models card; Home reminds you while this step is pending.

**THE 60-SECOND WALKTHROUGH**

Create a project → Upload an image set → Watch the task bar → Annotate five dimensions → Run text and vision analyses → Export and Methods paragraph